

Onlicorne: optimal transportation distances from sample streams

Anonymous Authors¹

Abstract

Optimal Transport (OT) distances are now routinely used as loss functions in ML tasks. Yet, computing OT distances between arbitrary (i.e. not necessarily discrete) probability distributions remains an open problem. This paper introduces a new online estimator of entropy-regularized OT distances between two such arbitrary distributions. It uses streams of samples from both distributions to iteratively enrich a non-parametric representation of the transportation plan. Compared to the classic Sinkhorn algorithm, our method leverages new samples at each iteration, which enables a consistent estimation of the true regularized OT distance. We cast our algorithm as a block-convex mirror descent in the space of positive distributions, which enables a theoretical analysis of its convergence. We numerically illustrate the performance of our method in comparison with concurrent approaches.

1. Introduction

Optimal transport (OT) distances are fundamental in statistical learning, both as a tool for analyzing the convergence of various algorithms (Canas & Rosasco, 2012; Dalalyan & Karagulyan, 2019), and as a data-dependent term for tasks as diverse as supervised learning (Frogner et al., 2015), unsupervised generative modeling (Martin Arjovsky, 2017) or domain adaptation (Courty et al., 2016). OT lifts a given distance over data points living in space \mathcal{X} into a distance on the space $\mathcal{P}(\mathcal{X})$ of probability distributions over this data space \mathcal{X} . We refer to the monograph (Santambrogio, 2015) for a detailed mathematical treatment. This distance has many favorable geometrical properties. In particular it allows one to compare distributions having disjoint supports. Computing OT distance is usually performed by sampling once from the input distributions and solving a discrete

linear program (LP), due to Kantorovich (1942). This approach is numerically costly and statistically inefficient (Weed et al., 2019). The optimisation problem depends on a fixed sampling of points from the data. It is therefore not adapted to machine learning setting where data is resampled continuously (e.g. in GANs), or accessed in an online manner. The goal of this paper is to develop an efficient online method able to estimate OT distances between continuous distributions — we will use a stream of data to refine the optimal transport solution. For this, we will adapt the celebrated Sinkhorn algorithm to an online setting.

Regularized OT. To alleviate both the computational and statistical burdens of OT, it is common to regularize the Kantorovich LP. The most successful approach in this direction is to use an entropic barrier penalty. When dealing with discrete distributions, this yields a problem that can be solved numerically using Sinkhorn-Knopp’s matrix balancing algorithm (Sinkhorn, 1964; Sinkhorn & Knopp, 1967). This approach was pushed forward for ML applications by Cuturi (2013). Sinkhorn distances are smooth and amenable to GPU computations, which makes them suitable as a loss function in model training. Sinkhorn distances are ε -accurate approximation of OT in $O(n^2/\varepsilon^3)$ for a number n of samples (Altschuler et al., 2017) (in contrast to the $O(n^3)$ complexity for an exact solution). Moreover, the optimal value of the regularized problem does not suffer from the curse of dimensionality (Genevay et al., 2019), since the average error using n random samples decay like $O(\varepsilon^{-d/2}/\sqrt{n})$, in sharp contrast with the slow $O(1/n^{1/d})$ error decay of OT (Weed et al., 2019). The regularized value can be sharpened by entropic debiasing (Feydy et al., 2019).

Handling streams of samples. The Sinkhorn algorithm operates in two distinct phases: draw samples from the distributions and evaluate a pairwise distance matrix in the first phase ; balance this distance matrix using Sinkhorn-Knopp iterations in the second phase. This two-step approach does not estimate the true regularized OT distance, and cannot handle samples provided as a stream, e.g. renewed at each training iteration of an outer algorithm. A cheap fix is to use Sinkhorn over mini-batches (see for instance Genevay et al. (2018) for an application to GANs). Yet this introduces a strong estimation bias, especially in high dimension —see Fatras et al. (2019) for a mathematical analysis.

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Continuous OT. Extending OT computations to arbitrary distributions (possibly having continuous densities) without relying on a fixed a priori sampling is an emerging topic of interest. A special case is the semi-discrete setting, where one of the two distributions is discrete. Without regularization, over an Euclidean space, this can be solved efficiently using the computation of Voronoi-like diagrams (Mérigot, 2011). This idea can be extended to entropic-regularized OT (Cuturi & Peyré, 2018), and can also be coupled with stochastic optimization method (Genevay et al., 2016) to tackle high dimensional problems (see also Staib et al. (2017) for an extension to Wasserstein barycenters).

When dealing with arbitrary continuous densities, which are accessed through a stream of random sample, the challenge is to approximate the (continuous) dual variables of the regularized Kantorovich LP using parametric or non-parametric functions. For application to generative model fitting, one can use deep networks, which leads to alternative formulation to Generative Adversarial Networks (GANs) (Martin Arjovsky, 2017) (see also Seguy et al. (2018) for an extension to the estimation of transportation maps). There is however no theoretical guarantees for this type of dual approximations, due to the non-convexity of the resulting optimization problem. The only mathematically rigorous approach in this direction is to use reproducing Hilbert space representations of potentials (Genevay et al., 2016). However, this approach relies on generic functional approximations, that results in very slow convergence.

Contribution. Our paper proposes a new take on continuous OT estimation. We construct a non-parametric representation of the potentials by extending the discrete Sinkhorn algorithm to the continuous setting.

2. Background: optimal transport distances

We recall the definition of optimal transport distances between arbitrary distributions (i.e. not necessarily discrete), then review how these are estimated using finite samples.

2.1. Optimal transport distances and algorithms

Wasserstein distances. We consider a complete metric space (\mathcal{X}, d) (assumed to be compact for simplicity), equipped with a continuous cost function $C(x, y) \in \mathbb{R}$ for any $(x, y) \in \mathcal{X}^2$ (assumed to be symmetric also for simplicity). Optimal transport lifts this *ground cost* into a cost between probability distributions over the space \mathcal{X} .

The Wasserstein cost between two probability distribution $(\alpha, \beta) \in \mathcal{P}(\mathcal{X})$ is defined as the minimal cost required to move each element of mass of α to each element of mass of β . It rewrites as the solution of a linear problem (LP) over the set of transportation plans (which are probability

distribution π over $\mathcal{X} \times \mathcal{X}$):

$$\mathcal{W}_C(\alpha, \beta) \triangleq \min_{\pi \in \mathcal{P}(\mathcal{X}^2)} \{ \langle C, \pi \rangle : \pi_1 = \alpha, \pi_2 = \beta \}, \quad (1)$$

where we denote $\langle C, \pi \rangle \triangleq \int C(x, y) d\pi(x, y)$. Here, $\pi_1 = \int_{y \in \mathcal{X}} d\pi(\cdot, y)$ and $\pi_2 = \int_{x \in \mathcal{X}} d\pi(x, \cdot)$ are the first and second marginals of the transportation plan π . When $C = d^p(x, y)$ is the p^{th} power of the ground distance, with $p \geq 1$, then $\mathcal{W}_C^{\frac{1}{p}}$ is itself a distance over $\mathcal{P}(\mathcal{X})$, whose associated topology is the one of the convergence in law (Santambrogio, 2015).

Entropic regularization and Sinkhorn algorithm. The solutions of (2) can be approximated by a strictly convex optimisation problem, where an entropic term is added to the linear objective to force curvature. The so-called Sinkhorn cost is then

$$\mathcal{W}_{C, \varepsilon}(\alpha, \beta) \triangleq \min_{\substack{\pi \in \mathcal{P}(\mathcal{X} \times \mathcal{X}) \\ \pi_1 = \alpha, \pi_2 = \beta}} \langle C, \pi \rangle + \varepsilon \text{KL}(\pi | \alpha \otimes \beta), \quad (2)$$

where the Kulback-Leibler divergence is defined as $\text{KL}(\pi | \alpha \otimes \beta) \triangleq \int \log(\frac{d\pi}{d\alpha d\beta}) d\pi$ (which is thus equal to the mutual information of π). This can be shown to approximate up to an $\varepsilon \log(\varepsilon)$ error $\mathcal{W}_C(\alpha, \beta)$ (see (Genevay et al., 2019)), which is recovered in the limit $\varepsilon = 0$. The regularized problem (2) admits a dual form, which is a maximization problem over the space of continuous *potential* function:

$$\max_{f, g \in \mathcal{C}(\mathcal{X})} \langle f, \alpha \rangle + \langle g, \beta \rangle - \varepsilon \langle \alpha \otimes \beta, \exp(\frac{f \oplus g - C}{\varepsilon}) \rangle + \varepsilon, \quad (3)$$

where $\langle f, \alpha \rangle \triangleq \int f(x) d\alpha(x)$ and $(f \oplus g - C)(x) \triangleq f(x) + g(y) - C(x, y)$. The major interest of problem (3) is that it can be solved by alternated maximization, which can be done in closed form. At iteration t , the iterations are simply

$$f_{t+1}(\cdot) = -T_{C, \varepsilon}(g_t, \beta), \quad g_{t+1}(\cdot) = -T_{C, \varepsilon}(f_{t+1}, \alpha),$$

$$\text{with } T_{C, \varepsilon}(h, \mu) \triangleq \int_{y \in \mathcal{X}} \exp(\frac{h(y) - C(\cdot, y)}{\varepsilon}) d\mu(y). \quad (4)$$

The operation $h \mapsto T_{C, \varepsilon}(h, \mu)$ maps a continuous function to another continuous function, and is a smooth approximation of the celebrated C -transform of OT (Santambrogio, 2015). We thus refer to it as a *soft C-transform*. The notation $f_t(\cdot)$ emphasizes the fact that f_t and g_t are continuous functions.

It can be shown that $(f_t)_t$ and $(g_t)_t$ converge in $(\mathcal{C}(\mathcal{X}), \|\cdot\|_{\text{var}})$ to a solution (f^*, g^*) of (3), where $\|f\|_{\text{var}} = \max_x f(x) - \min_x f(x)$ is the so-called variation norm. This is true for any initialisation f_0 . Convergence is due to the strict contraction of the operators $T_C(\cdot, \beta)$ and $T_C(\cdot, \alpha)$ in the space $(\mathcal{C}(\mathcal{X}), \|\cdot\|_{\text{var}})$ (Peyré et al., 2019).

2.2. Estimating OT distances with realizations

Iterations (4) cannot be implemented when dealing with generic distributions (α, β) , because it involves continuous functions (f_t, g_t) . When the input distribution are discrete (or equivalently that \mathcal{X} is a finite set) then these function can be stored using discrete sample, and algorithm (4) is equivalent to the celebrated Sinkhorn's algorithm (Sinkhorn, 1964; Sinkhorn & Knopp, 1967), which is often implemented over the scaling variable $(e^{f_t/\varepsilon}, e^{g_t/\varepsilon})$. More precisely, when dealing with empirical distributions $\hat{\alpha} = \sum_{i=1}^n a_i \delta_{x_i}$ and $\hat{\beta} = \sum_{i=1}^n a_i \delta_{y_i}$, then the approximation of $\mathcal{W}_{C, \varepsilon}(\alpha, \beta)$ using Sinkhorn iterations (4) compute $u_t \triangleq (e^{f_t(x_i)/\varepsilon})_{i=1}^n, v_t \triangleq (e^{g_t(y_i)/\varepsilon})_{i=1}^n$ as

$$u_{t+1} = \frac{1}{K(v_t \odot a)} \quad \text{and} \quad v_{t+1} = \frac{1}{K^\top(u_{t+1} \odot b)}$$

where $K = (e^{-\frac{C(x_i, y_j)}{\varepsilon}})_{i,j=1}^n \in \mathbb{R}^{n \times n}$. The algorithm thus operates in two phases: a first one, during which the kernel matrix K is computed (with a cost in $O(n^2 d)$, where d is the dimension of \mathcal{X}), and a second one, during which it is balanced (each iteration having a cost in $O(n^2)$).

The goal of this paper is to go beyond this discrete setting, and handle generic distributions (possibly having continuous densities). In particular, our numerical scheme manipulates continuous functions though an adapted parameterization which is automatically refined during the iterations.

2.3. Consistency and bias

Sample complexity results. The OT and Sinkhorn distances between two arbitrary distributions α and β can be approximated by the distance between discrete realizations $\hat{\alpha}_n = \frac{1}{n} \sum_i \delta_{x_i}, \hat{\beta}_n = \frac{1}{n} \sum_i \delta_{y_i}$, where $(x_i)_i$ and $(y_i)_i$ are i.i.d samples from α and β . Consistency holds, as $\mathcal{W}_{C, (\varepsilon)}(\hat{\alpha}_n, \hat{\beta}_n) \rightarrow \mathcal{W}_{C, (\varepsilon)}(\alpha, \beta)$, with a convergence rate in $\mathcal{O}(n^{-1/2})$ for Sinkhorn distances and $\mathcal{O}(n^{-1/d})$ for Wasserstein distances.

Bias in distance estimation. Although consistency is a reassuring result, the sample complexity of transport in high dimensions with low regularization remains high. For computational reasons, we cannot choose n to be much more than 10^5 . We may wonder whether we can improve the estimation of $\mathcal{W}_C(\alpha, \beta)$ using several sets of samples $(\hat{\alpha}_n^t)_t$ and $(\hat{\beta}_n^t)_t$. Those should be of reasonable size to allow Sinkhorn estimation, and may for example come from a temporal stream. (Genevay et al., 2018) proposes to use the Monte-Carlo estimate $\hat{\mathcal{W}}(\alpha, \beta) = \frac{1}{T} \sum_{t=1}^T \mathcal{W}(\hat{\alpha}_n^t, \hat{\beta}_n^t)$. However, this yields a wrong estimation as the distance $\mathcal{W}(\hat{\alpha}_n, \hat{\beta}_n)$ between discrete realizations is a *biased* estimator of $\mathcal{W}(\alpha, \beta)$:

$$\mathcal{W}(\alpha, \beta) \neq \mathbb{E}_{\hat{\alpha}_n \sim \alpha, \hat{\beta}_n \sim \beta}[\mathcal{W}(\hat{\alpha}_n, \hat{\beta}_n)].$$

Bias in gradients. In several applications, the distance $\mathcal{W}(\alpha, \beta)$ is used as a loss function. This is the case in generative modeling, when we parametrize α as the push-forward of some noise distribution μ through a neural network g_θ . We are then interested in computing the displacement gradient $\delta_\alpha \mathcal{W}(\alpha, \beta) \in \mathcal{P}(\mathcal{X})$, in order to train θ by backpropagation. This gradient turns out to be the spatial derivative $\nabla_x f^*$ of the solution of (3). Yet, similarly, estimating this gradient through sampling is also biased, as $f^*(\alpha, \beta) \neq \mathbb{E}_{\hat{\alpha}_n \sim \alpha, \hat{\beta}_n \sim \beta}[f^*(\hat{\alpha}_n, \hat{\beta}_n)]$. Our approach should help reducing this bias.

3. OT distances from sample streams

We introduce an online adaptation of the Sinkhorn algorithm in this section. We wish to construct an estimator of $\mathcal{W}(\alpha, \beta)$ from multiple sets of samples $(\hat{\alpha}^t)_t = \frac{1}{n} \sum_{i=1}^n \delta_{x_i^t}$ and $(\hat{\beta}^t)_t$. This estimator should successively use these samples to enrich a representation of the solution of (3), that may be arbitrary complex. $(\hat{\alpha}_n^t)_t$ and $(\hat{\beta}_n^t)_t$ may be seen as mini-batches within a training procedure, or as a temporal stream. We first introduce the intuitions behind the construction of our algorithm, before casting it as a non-convex stochastic mirror descent problem.

3.1. Online Sinkhorn iterations

From (3), along the Sinkhorn optimisation trajectory, the potential f_t is always the negative logarithm an infinite mixture of kernel functions $\kappa_y : x \rightarrow \exp(-\frac{C(\cdot, y)}{\varepsilon})$:

$$\exp(-\frac{f_t(\cdot)}{\varepsilon}) = \int_{y \in \mathcal{Y}} \exp(g_t(y)) \exp(-\frac{C(\cdot, y)}{\varepsilon}) d\beta(y),$$

and similarly for g_t . Our algorithm will construct a sequence (\hat{f}_t, \hat{g}_t) that behaves like g_t and f_t in the long run. The strong structural property of the continuous potentials suggests to express $\exp(-\frac{f_t(\cdot)}{\varepsilon})$ as a finite mixture of kernel functions. That is, \hat{f}_t and \hat{g}_t are continuous functions constructed from the weights $(p_i^s)_{i,s}, (q_i^s)_{i,s} > 0$ and positions $(y_i^s)_{i,t} \in \mathcal{Y}, (x_i^s)_{i,s \leq t} \in \mathcal{X}$:

$$\begin{aligned} \hat{f}_t(\cdot) &= -\log \sum_{s=1}^t \sum_{i=1}^{n_t} \exp(\frac{p_i^s - C(\cdot, y_i^s)}{\varepsilon}) \\ \hat{g}_t(\cdot) &= -\log \sum_{i=1}^{n_t} \exp(\frac{q_i^s - C(x_i^s, \cdot)}{\varepsilon}). \end{aligned}$$

Provided with fresh samples $(x_i^t)_i$ and $(y_i^t)_i$, a naive approach updates the potentials using a noisy soft C-transform:

$$\hat{f}_{t+1} = T(g_t, \hat{\beta}_t), \quad g_{t+1} = T(\hat{f}_{t+1}, \hat{\alpha}_t),$$

which translates into setting all $(p_i^s)_{s \leq 0}$ to 0, and to set each weight $p_i^{t+1} = T(g_t, \hat{\beta}_t)(x_i^{t+1})/n_t$.

As the variance of the updates does not reduce, this algorithm does not converge. It is however possible to show, through contractance of the random operator $T(\cdot, \hat{\beta})$, that the Markov chain (f_t, g_t) that it defines follows at the limit a stationary distribution that does not depend on initialisation.

The variance introduced must therefore be reduced in some way—in other words, we cannot afford to forget past iterates to obtain a consistent estimation of potentials. We introduce a learning rate in Sinkhorn iterations, that averages the past representations and the newly computed noisy C-transforms.

$$\begin{aligned} \exp(-\hat{f}_{t+1}) &= (1 - \eta_t) \exp(-\hat{f}_t) + \eta_t T(g_t, \hat{\beta}_t), \\ \exp(-\hat{g}_{t+1}) &= (1 - \eta_t) \exp(-\hat{g}_t) + \eta_t T(f_t, \hat{\alpha}_t). \end{aligned}$$

Performing the averaging in the space of $\exp(-\hat{f}_t)$ yields simple updates for the weights $(p_i^s)_s$, and is crucial for our theoretical understanding. In essence, the weights of past samples are reduced by a constant factor, while new weights are computed from the soft c-transform of $g_t(\cdot)$ at points y_i^t . Note that we perform *simultaneous updates* of f_t and g_t , once again from theoretical considerations.

3.2. Algorithm and complexity

We provide the pseudo code for online Sinkhorn in Alg. 1. The iteration t has a complexity in $\mathcal{O}(t n_t^2)$, due to the evaluation of the distances $C(x_i^t, y_i^s)$ for all $s < t$, and to the computation of the C-transforms. Online Sinkhorn constructs the distance matrix C on the fly, parallel to updating the potentials f and g . In total, its computation cost after seeing n samples is $\mathcal{O}(n^2)$, and its memory cost is $\mathcal{O}(n)$.

Refinements. Optionally, we may also compute out-of-loop averages of potentials

$$\begin{aligned} \exp(-\bar{f}_{t+1}) &= (1 - \gamma_t) \exp(-\bar{f}_t) + \gamma_t \exp(-\hat{f}_t), \\ \exp(-\bar{g}_{t+1}) &= (1 - \gamma_t) \exp(-\bar{g}_t) + \gamma_t \exp(-\hat{g}_t), \end{aligned}$$

to further reduce the estimation variance for an unchanged complexity.

Algorithm 1 Online Sinkhorn

Input: Distribution α and β , ε , η_t , γ_t
 Initialize $\mathcal{P} = []$, $\mathcal{Q} = []$, $\mathcal{X} = []$, $\mathcal{Y} = []$
for $t = 0, \dots, T$ **do**
 for $p \in \mathcal{P}$, $q \in \mathcal{Q}$ **do**
 $p \leftarrow p + \varepsilon \log(1 - \eta_t)$,
 $q \leftarrow q + \varepsilon \log(1 - \eta_t)$
 Sample $(x_i^t)_{i=1,n} \sim \alpha$, $(y_i^t)_{i=1,n} \sim \beta$
 for $y \in y_1^t, \dots, y_n^t$, $x \in x_1^t, \dots, x_n^t$ **do**
 $p_i^t \leftarrow -\varepsilon \log \frac{\eta_t}{n_t} \sum_{p \in \mathcal{P}, y \in \mathcal{Y}} \exp(\frac{p - C(x_i^t, y)}{\varepsilon})$
 $q_i^t \leftarrow -\varepsilon \log \frac{\eta_t}{n_t} \sum_{p \in \mathcal{P}, x \in \mathcal{X}} \exp(\frac{q - C(x, y_i^t)}{\varepsilon})$
 Add $(x_i^t)_i$ to \mathcal{X} , $(y_i^t)_i$ to \mathcal{Y}
 Add $(q_i^t)_i$ to \mathcal{Q} , $(p_i^t)_i$ to \mathcal{P}
 Returns representations of $f_T : (\mathcal{Y}, \mathcal{P})$, $g_T : (\mathcal{X}, \mathcal{Q})$

References

- Altschuler, J., Niles-Weed, J., and Rigollet, P. Near-linear time approximation algorithms for optimal transport via sinkhorn iteration. In *Advances in Neural Information Processing Systems*, pp. 1964–1974, 2017.
- Canas, G. and Rosasco, L. Learning probability measures with respect to optimal transport metrics. In *Advances in Neural Information Processing Systems*, pp. 2492–2500, 2012.
- Courty, N., Flamary, R., Tuia, D., and Rakotomamonjy, A. Optimal transport for domain adaptation. *IEEE transactions on pattern analysis and machine intelligence*, 39(9): 1853–1865, 2016.
- Cuturi, M. Sinkhorn distances: Lightspeed computation of optimal transport. In *Advances in Neural Information Processing Systems* 26, pp. 2292–2300, 2013.
- Cuturi, M. and Peyré, G. Semidual regularized optimal transport. *SIAM Review*, 60(4):941–965, 2018. doi: 10.1137/18M1208654. URL <https://arxiv.org/abs/1811.05527>.
- Dalalyan, A. S. and Karagulyan, A. User-friendly guarantees for the langevin monte carlo with inaccurate gradient. *Stochastic Processes and their Applications*, 129(12): 5278–5311, 2019.
- Fatras, K., Zine, Y., Flamary, R., Gribonval, R., and Courty, N. Learning with minibatch wasserstein: asymptotic and gradient properties. *arXiv preprint arXiv:1910.04091*, 2019.
- Feydy, J., Séjourné, T., Vialard, F.-X., Amari, S., Trounev, A., and Peyré, G. Interpolating between optimal transport and mmd using sinkhorn divergences. In *Proc. AIS-*

- TATS'19, 2019. URL <https://arxiv.org/abs/1810.08278>.
- Frogner, C., Zhang, C., Mobahi, H., Araya, M., and Poggio, T. A. Learning with a wasserstein loss. In *Advances in Neural Information Processing Systems*, pp. 2053–2061, 2015.
- Genevay, A., Cuturi, M., Peyré, G., and Bach, F. Stochastic optimization for large-scale optimal transport. In Lee, D. D., Luxburg, U. V., Guyon, I., and Garnett, R. (eds.), *Proc. NIPS'16*, pp. 3432–3440. Curran Associates, Inc., 2016. URL <http://hal.archives-ouvertes.fr/hal-01321664>.
- Genevay, A., Peyré, G., and Cuturi, M. Learning generative models with sinkhorn divergences. In *Proc. AISTATS'18*, pp. 1608–1617, 2018. URL <https://arxiv.org/abs/1706.00292>.
- Genevay, A., Chizat, L., Bach, F., Cuturi, M., and Peyré, G. Sample complexity of sinkhorn divergences. In *Proc. AISTATS'19*, 2019. URL <https://arxiv.org/abs/1810.02733>.
- Kantorovich, L. On the transfer of masses (in russian). *Doklady Akademii Nauk*, 37(2):227–229, 1942.
- Martin Arjovsky, Soumith Chintala, L. B. Wasserstein generative adversarial network. *Proc ICML'17*, 60(4):941–965, 2017.
- Mérogot, Q. A multiscale approach to optimal transport. In *Computer Graphics Forum*, volume 30, pp. 1583–1592. Wiley Online Library, 2011.
- Peyré, G., Cuturi, M., et al. Computational optimal transport. *Foundations and Trends® in Machine Learning*, 11(5-6): 355–607, 2019.
- Santambrogio, F. Optimal transport for applied mathematicians. *Birkhäuser, NY*, 55(58-63):94, 2015.
- Seguy, V., Damodaran, B. B., Flamary, R., Courty, N., Rolet, A., and Blondel, M. Large-scale optimal transport and mapping estimation. *Proc. ICLR*, 2018.
- Sinkhorn, R. A relationship between arbitrary positive matrices and doubly stochastic matrices. *Ann. Math. Statist.*, 35:876–879, 1964.
- Sinkhorn, R. and Knopp, P. Concerning nonnegative matrices and doubly stochastic matrices. *Pacific Journal of Mathematics*, 21(2):343–348, 1967.
- Staib, M., Claici, S., Solomon, J. M., and Jegelka, S. Parallel streaming wasserstein barycenters. In *Advances in Neural Information Processing Systems*, pp. 2647–2658, 2017.
- Weed, J., Bach, F., et al. Sharp asymptotic and finite-sample rates of convergence of empirical measures in wasserstein distance. *Bernoulli*, 25(4A):2620–2648, 2019.