

A Proofs

We first introduce two useful known lemmas, and prove the propositions in their order of appearance.

A.1 Useful lemmas

First, under [Assumption 1](#), we note that the soft C -transforms are uniformly contracting on the distribution space $\mathcal{P}(\mathcal{X})$. This is clarified in the following lemma, extracted from Vialard [\[43\]](#), Proposition 19. We refer the reader to the original references for proofs.

Lemma 1. *Under [Assumption 1](#), let $\kappa = 1 - \exp(-L\text{diam}(\mathcal{X}))$. For all $\hat{\alpha} \in \mathcal{P}(\mathcal{X})$ and $\hat{\beta} \in \mathcal{P}(\mathcal{X})$, for all $f, f', g, g' \in \mathcal{C}(\mathcal{X})$,*

$$\|T_{\hat{\alpha}}(f') - T_{\hat{\alpha}}(f)\|_{\text{var}} \leq \kappa \|f - f'\|_{\text{var}}, \quad \|T_{\hat{\beta}}(g) - T_{\hat{\beta}}(g')\|_{\text{var}} \leq \kappa \|g - g'\|_{\text{var}}.$$

We will also need a uniform law of large numbers for functions. The following lemma is a consequence of Example 19.7 and Lemma 19.36 of Van der Vaart [\[42\]](#), and is copied in Lemma B.6 in Mairal [\[29\]](#).

Lemma 2. *Under [Assumption 1](#), let $(f_t)_t$ be an i.i.d sequence in $\mathcal{C}(\mathcal{X})$, such that $\mathbb{E}[f_0] = f \in \mathcal{C}(\mathcal{X})$. Then there exists $A > 0$ such that, for all $n > 0$,*

$$\mathbb{E} \sup_{x \in \mathcal{X}} \left| \frac{1}{n} \sum_{i=1}^n f_i(x) - f(x) \right| \leq \frac{A}{\sqrt{n}}.$$

Finally, we need a result on running averages using the sequence $(\eta_t)_t$. The following result stems from a simple Abel transform of the law of large number, and is established by Mairal [\[29\]](#), Lemma B.7.

Lemma 3. *Let $(\eta_t)_t$ be a sequence of weights meeting [Assumption 2](#). Let $(X_t)_t$ be an i.i.d sequence of real-valued random variables with existing first moment $\mathbb{E}[X_0]$. We consider the sequence $(\bar{X}_t)_t$ defined by $\bar{X}_0 \triangleq X_0$ and*

$$\bar{X}_t \triangleq (1 - \eta_t) \bar{X}_{t-1} + \eta_t X_t.$$

Then $\bar{X}_t \rightarrow_{t \rightarrow \infty} \mathbb{E}[X_0]$.

A.2 Proof of [Proposition 1](#)

Proof. We use Theorem 1 from Diaconis and Freedman [\[14\]](#). For this, we simply note that the space $\mathcal{C}(\mathcal{X}) \times \mathcal{C}(\mathcal{X})$ in which the chain $x_t \triangleq (f_t, g_t)_t$, endowed with the metric $\rho((f_1, g_1), (f_2, g_2)) = \|f_1 - f_2\|_{\text{var}} + \|g_1 - g_2\|_{\text{var}}$, is complete and separable (the countable set of polynomial functions are dense in this space, for example). We consider the operator $A_\theta \triangleq T_{\hat{\beta}}(T_{\hat{\alpha}}(\cdot))$. $\theta \triangleq (\hat{\alpha}, \hat{\beta})$ denotes the random variable that is sampled at each iteration. We have the following recursion:

$$x_{t+2} = A_{\theta_t}(x_t).$$

From [Lemma 1](#), for all $\hat{\alpha} \in \mathcal{P}(\mathcal{X})$, $\hat{\beta} \in \mathcal{P}(\mathcal{X})$, A_θ with $\theta = (\hat{\alpha}, \hat{\beta})$ is contracting, with module $\kappa_\theta < \kappa < 1$. Therefore

$$\int_{\theta} \kappa_\theta d\mu(\theta) < 1, \quad \int_{\theta} \log \kappa_\theta d\mu(\theta) < 0.$$

Finally, we note, for all $f \in \mathcal{C}(\mathcal{X})$

$$\|T_{\hat{\beta}}(T_{\hat{\alpha}}(f))\|_{\infty} \leq \|f\|_{\infty} + 2 \max_{x, y \in \mathcal{X}} C(x, y),$$

therefore $\rho(A_\theta(x_0), x_0) \leq 2\|x_0\|_{\infty} + 2 \max_{x, y \in \mathcal{X}} C(x, y)$ for all $\theta = (\hat{\alpha}, \hat{\beta})$. The regularity condition of the theorem are therefore met. Each of the induced Markov chains $(f_{2t}, g_{2t})_t$ and $(f_{2t+1}, g_{2t+1})_t$ has a unique stationary distribution. These stationary distributions are the same: the stationary distribution is independent of the initialisation and both sequences differs only by their initialisation. Therefore $(f_t, g_t)_t$ have a unique stationary distribution (F_{∞}, G_{∞}) . \square

448 A.3 Proof of Proposition 2

449 For presentation purpose, we first show that the “slowed-down” online Sinkhorn algorithm converges
450 in the absence of noise. We then turn to prove Proposition 2.

451 A.3.1 Noise-free online Sinkhorn

452 **Proposition 5.** *We suppose that $\hat{\alpha}_t = \alpha$, $\hat{\beta}_t = \beta$ for all t . Then the updates (6) yields a (deterministic)
453 sequence $(f_t, g_t)_t$ such that*

$$\|\hat{f}_t - f^*\|_{\text{var}} + \|\hat{g}_t - g^*\|_{\text{var}} \rightarrow 0, \quad \frac{1}{2} \langle \alpha, f_t + T_\alpha(\hat{g}_t) \rangle + \langle \beta, \hat{g}_t + T_\beta(f_t) \rangle \rightarrow \mathcal{W}(\alpha, \beta).$$

454 Note that, as we perform *simultaneous* updates, we only obtain the convergence of $f_t \rightarrow f^* + A$, and
455 $g_t \rightarrow g^*$, where f^* and g^* are solutions of (1) and A is a constant depending on initialization.

456 The “slowed-down” Sinkhorn iterations converge toward an optimal potential couple, up to a constant
457 factor: this stems from the fact that we apply contractions in the space $(\mathcal{C}(\mathcal{X}), \|\cdot\|_{\text{var}})$ with a
458 contraction factor that decreases sufficiently slowly.

459 *Proof.* We write $(f_t, g_t)_t$ the sequence of iterates. Given a pair of optimal potentials (f^*, g^*) , we
460 write $u_t \triangleq f_t - f^*$, $v_t \triangleq g_t - g^*$, $u_t^T \triangleq T_\alpha(f_t) - g^*$ and $v_t^T \triangleq T_\beta(g_t) - f^*$. For all $t > 0$, we
461 observe that

$$\begin{aligned} \max u_{t+1} &= -\log \min \exp(-u_{t+1}) \\ &= -\log \left(\min \left((1 - \eta_t) \exp(-u_t) + \eta_t \exp(-v_t^T) \right) \right) \\ &\leq -\log \left((1 - \eta_t) \min \exp(-u_t) + \eta_t \min \exp(-v_t^T) \right) \\ &\leq -(1 - \eta_t) \log \min \exp(-u_t) - \eta_t \log \min \exp(-v_t^T) \\ &= (1 - \eta_t) \max u_t + \eta_t \max v_t^T, \end{aligned}$$

462 where we have used the algorithm recursion on the second line, $\min f + g \geq \min f + \min g$ on the
463 third line and Jensen inequality on the fourth line. Similarly

$$\min u_{t+1} \geq (1 - \eta_t) \min u_t + \eta_t \min v_t^T,$$

464 and mirror inequalities hold for v_t . Summing the four inequalities, we obtain

$$\begin{aligned} e_{t+1} &\triangleq \|u_{t+1}\|_{\text{var}} + \|v_{t+1}\|_{\text{var}} \\ &= \max u_{t+1} - \min u_{t+1} + \max v_{t+1} - \min v_{t+1} \\ &\leq (1 - \eta_t)(\|u_t\|_{\text{var}} + \|v_t\|_{\text{var}}) + \eta_t(\|u_t^T\|_{\text{var}} + \|v_t^T\|_{\text{var}}), \\ &\leq (1 - \eta_t)(\|u_t\|_{\text{var}} + \|v_t\|_{\text{var}}) + \eta_t \kappa (\|u_t\|_{\text{var}} + \|v_t\|_{\text{var}}), \end{aligned}$$

465 where we use the contractivity of the soft- C -transform, that guarantees that there exists $\kappa < 1$ such
466 that $\|v_t^T\|_{\text{var}} \leq \kappa \|v_t\|_{\text{var}}$ and $\|u_t^T\|_{\text{var}} \leq \kappa \|u_t\|_{\text{var}}$ [34].

467 Unrolling the recursion above, we obtain

$$\log e_t = \sum_{s=1}^t \log(1 - \eta_s(1 - \kappa)) + \log(e_0) \rightarrow -\infty,$$

468 provided that $\sum \eta_t = \infty$. The proposition follows. \square

469 *Proof of Proposition 2.* For discrete realizations $\hat{\alpha}$ and $\hat{\beta}$, we define the perturbation terms

$$\varepsilon_{\hat{\beta}}(\cdot) \triangleq f^* - T_{\hat{\beta}}(g^*), \quad \iota_{\hat{\alpha}}(\cdot) \triangleq g^* - T_{\hat{\alpha}}(f^*),$$

470 so that the updates can be rewritten as

$$\begin{aligned} \exp(-f_{t+1} + f^*) &= (1 - \eta_t) \exp(-f_t + f^*) + \eta_t \exp(-T_{\hat{\beta}_t}(g_t) + T_{\hat{\beta}_t}(g^*) + \varepsilon_{\hat{\beta}_t}) \\ \exp(-g_{t+1} + g^*) &= (1 - \eta_t) \exp(-g_t + g^*) + \eta_t \exp(-T_{\hat{\alpha}_t}(f_t) + T_{\hat{\alpha}_t}(f^*) + \iota_{\hat{\alpha}_t}). \end{aligned}$$

471 We denote $u_t \triangleq -f_t + f^*$, $v_t \triangleq -g_t + g^*$, $u_t^T \triangleq T_{\hat{\beta}_t}(f_t) - T_{\hat{\beta}_t}(f^*)$, $v_t^T \triangleq T_{\hat{\beta}_t}(g_t) - T_{\hat{\beta}_t}(g^*)$.
 472 Reusing the same derivations as in the proof of [Proposition 5](#), we obtain

$$\begin{aligned} \|u_{t+1}\|_{\text{var}} &\leq (1 - \eta_t) \|u_t\|_{\text{var}} \\ &\quad + \eta_t \log \left(\max_{x, y \in \mathcal{X}} \exp(\varepsilon_{\hat{\beta}_t}(x) - \varepsilon_{\hat{\beta}_t}(y)) \exp(v_t^T(x) - v_t^T(y)) \right) \\ &\leq (1 - \eta_t) \|u_t\|_{\text{var}} + \eta_t \|v_t^T\|_{\text{var}} + \eta_t \|\varepsilon_{\hat{\beta}_t}\|_{\text{var}}, \end{aligned}$$

473 where we have used $\max_x f(x)g(x) \leq \max_x f(x) \max_x g(x)$ on the second line. Therefore, using
 474 the contractivity of the soft C -transform,

$$e_{t+1} \leq (1 - \tilde{\eta}_t) e_t + \frac{\tilde{\eta}_t}{1 - \kappa} (\|\varepsilon_{\hat{\beta}_t}\|_{\text{var}} + \|\iota_{\hat{\alpha}_t}\|_{\text{var}}), \quad (8)$$

475 where we set $e_t \triangleq \|u_t\|_{\text{var}} + \|v_t\|_{\text{var}}$, $\tilde{\eta}_t = \eta_t(1 - \kappa)$ and κ is set to be the biggest contraction
 476 factor over all empirical realizations $\hat{\alpha}_t, \hat{\beta}_t$ of the distributions α and β . It is upper bounded by
 477 $1 - e^{-L \text{diam}(\mathcal{X})}$, thanks to [Assumption 1](#) and [Lemma 1](#).

478 The realizations $\hat{\beta}_t$ and $\hat{\alpha}_t$ are sampled according to the same distribution $\hat{\alpha}$ and $\hat{\beta}$. We define the
 479 sequence r_t to be the running average of the variational norm of the (functional) error term:

$$r_{t+1} \triangleq (1 - \tilde{\eta}_t) r_t + \frac{\tilde{\eta}_t}{1 - \kappa} (\|\varepsilon_{\hat{\beta}_t}\|_{\text{var}} + \|\iota_{\hat{\alpha}_t}\|_{\text{var}}).$$

480 We thus have, for all $t > 0$, $e_t \leq r_t$. Using [Lemma 3](#), the sequence $(r_t)_t$ converges towards the scalar
 481 expected value

$$r_\infty \triangleq \frac{1}{1 - \kappa} \mathbb{E}_{\hat{\alpha}, \hat{\beta}} [\|\varepsilon_{\hat{\beta}}\|_{\text{var}} + \|\iota_{\hat{\alpha}}\|_{\text{var}}] > 0. \quad (9)$$

482 We now relate r_∞ to the number of samples n using a uniform law of large number result on
 483 parametric functions. We write $\hat{\beta} = \hat{\beta}_n$ to make explicit the dependency of the quantities on the
 484 batch size n .

485 Using [Lemma 2](#), we bound the quantity

$$\begin{aligned} E_n &\triangleq \mathbb{E}_{\hat{\beta}_n} \|\varepsilon_{\hat{\beta}_n}\|_{\text{var}} = \mathbb{E}_{\hat{\beta}_n} \|\exp(-T_{\hat{\beta}_n}(g_0^*)) - \exp(-T_{\hat{\beta}_n}(g_0^*))\|_\infty \\ &= \mathbb{E}_{Y_1, \dots, Y_n \sim \beta} \sup_{x \in \mathcal{X}} \left| \frac{1}{n} \sum_{i=1}^n \exp(g^*(Y_i)) - C(x, Y_i) \right| \\ &\quad - \mathbb{E}_{Y \sim \beta} [\exp(g_0^*(Y)) - C(x, Y)] \\ &= \mathbb{E} \sup_{x \in \mathcal{X}} \left| \frac{1}{n} \sum_{i=1}^n \varphi_i(x) - \varphi(x) \right|, \end{aligned}$$

486 where we have defined $\varphi_i : x \rightarrow \exp(g^*(Y_i)) - C(x, Y_i)$ and set φ to be the expected value of each
 487 φ_i . The compactness of \mathcal{X} ensures that the functions are square integrable and uniformly bounded.
 488 [Lemma 2](#) ensures that there exists $S(g^*)$ such that

$$E_n \leq \frac{S(g^*)}{\sqrt{n}}.$$

489 We now bound $\mathbb{E}_{\hat{\beta}_n} \|\varepsilon_{\hat{\beta}_n}\|_{\text{var}}$ using the quantity E_n . First, we observe that $\|\cdot\|_{\text{var}} = g_{\min}^* < g^* < 0$, and
 490 there exists $C_{\max} > 0$ such that $0 \leq C(x, y) \leq C_{\max}$ for all $x, y \in \mathcal{X}$, thanks to the [Assumption 1](#).

$$\begin{aligned} \delta &\triangleq \exp(-\|g^*\|_{\text{var}} - C_{\max}) \leq \exp(-T_{\hat{\beta}_n}(g^*)) \leq 1 \\ &\quad \exp(-\|g^*\|_{\text{var}} - C_{\max}) \leq \exp(-T_{\hat{\beta}_n}(g^*)) \leq 1, \end{aligned}$$

491 where we have used $g^* = \|g^*\|_{\text{var}}$. For all $x \in \mathcal{X}$,

$$|\varepsilon_{\hat{\beta}_n}| = \left| \log \frac{\exp(-T_{\hat{\beta}_n}(g^*))}{\exp(-T_{\hat{\beta}_n}(g^*))} \right| = \left| \log \left(1 + \frac{\exp(-T_{\hat{\beta}_n}(g^*)) - \exp(-T_{\hat{\beta}_n}(g^*))}{\exp(-T_{\hat{\beta}_n}(g^*))} \right) \right|. \quad (10)$$

492 We first obtain an upper-bound independent of n with the first equality in (10):

$$\|\varepsilon_{\hat{\beta}_n}\|_{\text{var}} \leq \|\varepsilon_{\hat{\beta}_n}\|_{\infty} \leq \|g^*\|_{\text{var}} + C_{\max}. \quad (11)$$

493 We now use the second expression in (10): for n large enough, $E_n < \delta$

$$\|\varepsilon_{\hat{\beta}_n}\|_{\text{var}} \leq \max(\log(1 + \frac{E_n}{\delta}), -\log(1 - \frac{E_n}{\delta})) = -\log(1 - \tilde{E}_n), \quad (12)$$

494 where we have set $\tilde{E}_n \triangleq \frac{E_n}{\delta}$. On the event $\Omega_n = \{\tilde{E}_n \leq \frac{1}{2}\}$, a simple calculation gives $-\log(1 - \tilde{E}_n) \leq (2 \log 2) \tilde{E}_n \leq 2\tilde{E}_n$. Thanks to Markov inequality, $\mathbb{P}[\tilde{E}_n > \frac{1}{2}] \leq 2\mathbb{E}[\tilde{E}_n]$. We then split the expectation over the event Ω_n , and use inequalities (12) and (11) on each conditional expectation:

$$\begin{aligned} \mathbb{E}\|\varepsilon_{\hat{\beta}_n}\|_{\text{var}} &= \mathbb{P}\left[\tilde{E}_n \leq \frac{1}{2}\right] \mathbb{E}\left[\|\varepsilon_{\hat{\beta}_n}\|_{\text{var}} \mid \tilde{E}_n \leq \frac{1}{2}\right] \\ &\quad + \mathbb{P}\left[\tilde{E}_n > \frac{1}{2}\right] \mathbb{E}\left[\|\varepsilon_{\hat{\beta}_n}\|_{\text{var}} \mid \tilde{E}_n > \frac{1}{2}\right] \\ &\leq \frac{2\varphi(\|g^*\|_{\text{var}} + C_{\max})S(g^*)}{\sqrt{n}} \\ &\leq \frac{4 \exp(\|g^*\|_{\text{var}} + C_{\max})S(g^*)}{\sqrt{n}} \triangleq \frac{A(g^*)}{\sqrt{n}} \end{aligned} \quad (13)$$

497 The constants S depends on the complexity of estimating the functional $x \rightarrow \int_y \exp(g^*(y) - C(x, y))d\beta(y)$ with samples from β . A parallel result holds for $\mathbb{E}_{\hat{\alpha}_n}\|\iota_{\hat{\alpha}_n}\|_{\text{var}}$. Therefore, there exists $A(f^*), A(g^*) > 0$ such that $r_{\infty} \leq \frac{A(f^*) + A(g^*)}{\sqrt{n}}$. As for all $t > 0$, $e_t \leq r_t \rightarrow_{t \rightarrow \infty} r_{\infty}$, the proposition follows, writing $A = A(f^*) + A(g^*)$.

501 The constant A is larger than $\exp(C_{\max})$ when $C_{\max} \rightarrow \infty$; Hence it behaves at least like $\exp(\frac{1}{\varepsilon})$ when $\varepsilon \rightarrow 0$.

503 Note that we have used twice a corollary of the law of large numbers: once when averaging over t with $t \rightarrow \infty$ (Eq. (9)), and once when averaging over n with n finite (Eq. (13)). \square

505 A.4 Proof of Proposition 3

506 In the proof of Proposition 2 and in particular Eq. (8), the term that prevents the convergence of e_t is

$$\eta_t(\|\varepsilon_{\hat{\beta}_t}\|_{\text{var}} + \|\iota_{\hat{\alpha}_t}\|_{\text{var}}),$$

507 which is not summable in general. We can control this term by increasing the size of $\hat{\alpha}_t$ and $\hat{\beta}_t$ with time, at a sufficient rate: this is what Assumption 3 ensures.

509 *Proof.* From Eq. (8), for all $t > 0$, we have

$$0 \leq e_{t+1} \leq (1 - \tilde{\eta}_t)e_t + \eta_t(\|\varepsilon_{\hat{\beta}_t}\|_{\text{var}} + \|\iota_{\hat{\alpha}_t}\|_{\text{var}}). \quad (14)$$

510 Taking the expectation and using the uniform law of large number (13),

$$\begin{aligned} \mathbb{E}e_{t+1} &\leq (1 - (1 - \kappa)\eta_t)\mathbb{E}e_t + \eta_t \frac{A}{\sqrt{n(t)}} \\ &= (1 - (1 - \kappa)\eta_t)\mathbb{E}e_t + A\eta_t w_t, \end{aligned} \quad (15)$$

511 where we have used the definition of $n(t)$ from Assumption 3 in the last line.

512 The proof follows from a simple asymptotic analysis of the sequence $(\mathbb{E}e_t)_t$, following recursion (15). For all $t > 0$,

$$\mathbb{E}e_{t+1} - \mathbb{E}e_t = -(1 - \kappa)\eta_t\mathbb{E}e_t + A\eta_t w_t \leq A\eta_t w_t \quad (16)$$

514 Therefore, from Assumption 3, $(\mathbb{E}e_{t+1} - \mathbb{E}e_t)_t$ is summable and $\mathbb{E}e_t \rightarrow_{t \rightarrow \infty} \ell \geq 0$. Let's assume $\ell > 0$. Summing (16) over t , we obtain

$$\mathbb{E}e_t \leq \mathbb{E}e_1 - (1 - \kappa) \sum_{s=1}^{t-1} \eta_s \mathbb{E}e_s + A \sum_{s=1}^{t-1} \eta_s w_s \rightarrow_{t \rightarrow \infty} -\infty,$$

516 which leads to a contradiction. Therefore $\mathbb{E}e_t \rightarrow_{t \rightarrow \infty} 0$. As $e_t \geq 0$ for all $t > 0$, this implies that $e_t \rightarrow_{t \rightarrow \infty} 0$ almost surely. \square

518 A.5 Proof of Proposition 4

519 *Proof.* The proof of Proposition 3 allows us to derive non-asymptotic rates for potential estimations
 520 using the online Sinkhorn algorithm. Let us set $\eta_t = \frac{\lambda}{t^a}$, $n(t) = \lceil Bt^{2b} \rceil$ in (14), so that Assumption 3
 521 is met. $\lceil \cdot \rceil$ denotes the ceiling function. We are left to study the recursion (15):

$$\delta_{t+1} \triangleq \mathbb{E}e_{t+1} \leq (1 - \frac{\lambda(1-\kappa)}{t^a})\delta_t + \frac{A\lambda}{\sqrt{B}t^{a+b}}$$

522 Following the derivations of Moulines and Bach [33, Theorem 2], we have the following bias-variance
 523 decomposed upper-bound, provided that $0 \leq a < 1$ and $a + b > 1$. For all $t > 0$,

$$\delta_t \leq (\delta_0 + \frac{AS}{(a+b-1)\sqrt{B}}) \exp(-\frac{S(1-\kappa)}{2}t^{1-a}) + \frac{2AS}{\sqrt{B}(1-\kappa)t^a}. \quad (17)$$

524 Let us now relate the iteration number t to the number of seen sample N . By definition

$$n_t = \sum_{s=1}^t n(s) \leq B \sum_{s=1}^t s^{2b} + t \leq t + \frac{(t+1)^{2b+1} - 1}{2b+1} \leq (2t)^{2b+1}.$$

525 Therefore, when we have seen N samples, the iteration number is superior to $t(N)$, and the expected
 526 error δ_N is of the order of $\delta_{t(N)}$, with

$$t(N) = (N/2)^{\frac{1}{2b+1}}. \quad (18)$$

527 We write $\delta_N = \delta_{t(N)}$. Replacing (18) in (17) yields

$$\delta_n \leq (\delta_0 + \frac{A\lambda}{(a+b-1)\sqrt{B}}) \exp\left(-\frac{\lambda(1-\kappa)}{2}(n/2)^{\frac{1-a}{2b+1}}\right) + \frac{2A\lambda}{\sqrt{B}(1-\kappa)(n/2)^{\frac{a}{2b+1}}}. \quad (19)$$

528 We note that b and a should be as close to 0 as possible to reduce the bias term, while a should be
 529 as close to 1 and b as close to 0 as possible to reduce the variance term. Of course, b should remain
 530 larger than $1 - a$ to ensure convergence.

531 To obtain the best asymptotical rates (the error is always dominated by the variance term), we set
 532 $a = 1 - \iota$, $b = 2\iota$, with $\iota \gtrsim 0$. This yields

$$\begin{aligned} \delta_n &\leq (\delta_0 + \frac{A\lambda}{\iota\sqrt{B}}) \exp\left(-\frac{\lambda(1-\kappa)}{2}(n/2)^{\frac{\iota}{1+4\iota}}\right) + \frac{2A\lambda}{\sqrt{B}(1-\kappa)(n/2)^{\frac{1-\iota}{1+4\iota}}} \\ &= \mathcal{O}(n^{-\frac{1-\iota}{1+4\iota}}). \end{aligned}$$

533 This rate is as close to the rate $\mathcal{O}(\frac{1}{n})$ as desired. We may then perform a last soft C -transform
 534 (using the n_t seen samples) over the estimated $f_{t(n)}, g_{t(n)}$ to obtain a estimated solution of the dual
 535 optimisation problem (2). The Sinkhorn potentials can therefore be estimated with *fast rates*. Note
 536 that the upper bound explodes when $\varepsilon \rightarrow 0$, as $C_{\max} \rightarrow \infty$, hence $A \rightarrow \infty$, and $(1 - \kappa) \rightarrow 0$. \square

537 **Estimating the Sinkhorn distance.** The Sinkhorn distance requires to estimate the integral

$$\mathcal{W}(\alpha, \beta) = \int_x f^*(x) d\alpha(x) + \int_y g^*(y) d\beta(y).$$

538 At iteration $t(n)$, with empirical realization $\bar{\alpha}_t$ and $\bar{\beta}_t$, containing n samples, we use the estimator

$$\hat{\mathcal{W}}(\alpha, \beta) = \frac{1}{n} \sum_{i=1}^n f_{t(n)}(x_i) + \frac{1}{n} \sum_{i=1}^n g_{t(n)}(y_i),$$

539 We can bound the estimation error $|\hat{\mathcal{W}}(\alpha, \beta) - \mathcal{W}(\alpha, \beta)| = \mathcal{O}(\frac{1}{\sqrt{n}})$, dominated by the integral
 540 evaluation noise. We thus recover a new estimator of the Sinkhorn distance with the same sample
 541 complexity as the batch Sinkhorn estimator [19]. Our estimator enjoys an original rate for estimating
 542 the potentials in $\|\cdot\|_{\text{var}}$.

Algorithm 2 Fully-corrective online Sinkhorn

Input: Distribution α and β , learning weights $(\eta_t)_t$ and batch-sizes $(n(t))_t$. **Set** $p_{i,1} = q_{i,1} = 0$ for $i \in (0, n_1]$
for $t = 0, \dots, T - 1$ **do**
 Sample $(x_i)_{(n_t, n_{t+1}]} \sim \alpha, (y_j)_{(n_t, n_{t+1}]} \sim \beta$.
 Evaluate $(\hat{f}_t(x_i))_{i=(0, n_{t+1}]}, (\hat{g}_t(y_j))_{j=(0, n_{t+1}]}$ using $(q_{i,t}, p_{i,t}, x_i, y_i)_{i=(0, n_t]}$ in (7).
 $q_{(0, n_{t+1}], t+1} \leftarrow \log \frac{1}{n} + (\hat{g}_t(y_i))_{(0, n_{t+1}]}, \quad p_{(n_t, n_{t+1}], t+1} \leftarrow \log \frac{1}{n} + (\hat{f}_t(x_i))_{(n_t, n_{t+1}]}$.
Returns: $\hat{f}_T : (q_{i,T}, y_i)_{(0, n_T]}$ and $\hat{g}_T : (p_{i,T}, x_i)_{(0, n_T]}$

Algorithm 3 Online Sinkhorn potentials in the discrete setting

Input: Distribution $\alpha \in \Delta^N$ and $\beta \in \Delta^N, x \in \mathbb{R}^{n \times d}, y \in \mathbb{R}^{n \times d}$, learning weights $(\eta_t)_t$
Set $p = q = -\infty \in \mathbb{R}^n$.
for $t = 1, \dots, T$ **do**
 $q \leftarrow q + \log(1 - \eta_t), p \leftarrow p + \log(1 - \eta_t)$.
 Sample $J_t \subset [1, N], I_t \subset [1, N]$ of size $n(t)$.
 for $i \in J_t$ **do**
 $q_i \leftarrow \log \left(\exp(q_i) + \exp \left(\log(\eta_t) - \log \frac{1}{N} \sum_{j=1}^N \exp(p_j - C(x_j, y_i)) \right) \right)$.
 for $i \in I_t$ **do**
 $p_i \leftarrow \log \left(\exp(q_i) + \exp \left(\log(\eta_t) - \log \frac{1}{M} \sum_{j=1}^M \exp(q_j - C(x_i, y_j)) \right) \right)$.
Returns $f_T : (q, y)$ and $g_T : (p, x)$

543 **B Online Sinkhorn variants**544 **B.1 Fully-corrective scheme**

545 We report the fully-corrective online Sinkhorn algorithm in [Algorithm 2](#). This algorithm also enjoys
546 almost sure convergence, provided that the following assumption is met.

547 **Assumption 4.** For all $t > 0$, the total batch-size $n_t = \frac{B}{w_t^2}$ is an integer. The step-size η_t and the
548 batch-size n_t grows so that $\sum w_t \eta_t < \infty$ and $\sum \eta_t = \infty$.

549 With full correction, the total number of observed samples n_t needs to grow at the same rate as the
550 single-iteration batch-size $n(t)$ in [Assumption 3](#). For $\eta_t = \frac{1}{t^a}, a \in (1/2, 1]$, it is sufficient to use a
551 constant batch-size $n(t) = B$ to meet [Assumption 4](#). We then have the following property

552 **Proposition 6.** Under [Assumption 1](#) and [4](#), the fully-corrective online Sinkhorn algorithm converges
553 almost surely:

$$\|\hat{f}_t - f^*\|_{\text{var}} + \|\hat{g}_t - g^*\|_{\text{var}} \rightarrow 0.$$

554 *Proof.* Using the fully-corrective scheme allows to replace $n(t)$ by $n_t = \sum_{s=0}^t n(s)$ in (15). The
555 proposition is then obtained in the same way as [Proposition 4](#). \square

556 **B.2 Online Sinkhorn for discrete distributions**

557 The online Sinkhorn algorithm takes a simpler form with discrete distributions. We derive it in
558 [Algorithm 3](#). We set α and β to have size N and M , respectively. We evaluate the potentials as

$$g_t(y) = -\log \sum_{j=1}^N \exp(p_j - C(x_j, y))$$

$$f_t(x) = -\log \sum_{j=1}^M \exp(q_j - C(x, y_j)),$$

Table 1: Schedules of batch-sizes and learning rates that ensures online Sinkhorn convergence.

Param. schedule	Online Sinkhorn	Fully-corrective online Sinkhorn
Batch size $n(t) = Bt^b$	$0 < b$	$0 \leq b$
Step size $\eta_t = \frac{1}{t^a}$	$a \geq 1 - \frac{b}{2}$	$\begin{cases} a > \frac{1}{2} - \frac{b}{2} & \text{and } b < 1 \\ a \geq 0 & \text{and } b \geq 1 \end{cases}$

where $(p_j)_{j \in [1, N]}$ and $(q_j)_{j \in [1, M]}$ are fixed-size vectors. Note that the computations written in Algorithm 3 are in log-space, as they should be implemented to prevent numerical overflows. The sets $|I|$ and $|J|$ can have varying sizes along the algorithm, which allows for example to speed-up the initial Sinkhorn iteration (§5.2). In this case, the cost matrix $\hat{C} = C(x_i, y_j)_{i, j}$ should be progressively recorded along the algorithm iterations.

B.3 Recapitulation on batch-sizes and learning rates

To provide practical guidance on choosing rates in batch-sizes $n(t)$ and step-sizes η_t , we can parametrize $\eta_t = \frac{1}{t^a}$ and $n(t) = Bt^b$ and study what is implied by Assumption 3 and Assumption 4. We summarize the schedules for which convergence is guaranteed in Table 1. Note that in practice, it is useful to replace t by $(1 + r t)$ in these schedules. We set $r = 0.1$ in all experiments.

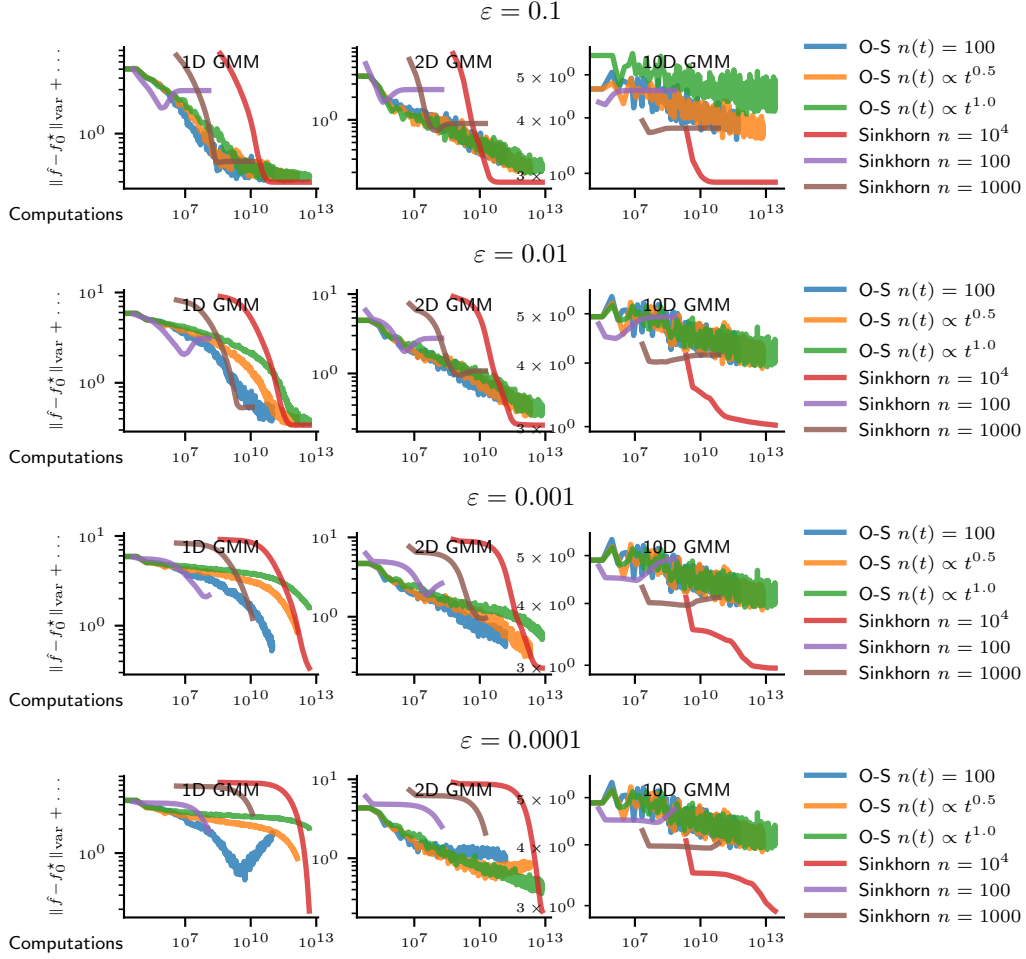


Figure 4: Performance of online Sinkhorn for various ε .

C Extra numerical experiments

We display and describe the supplementary figures mentioned in the main text, as well as experimental details useful for reproduction.

C.1 Online Sinkhorn and variants

Grids and details for §5.1. We set $(\eta_t, n(t)) = (\frac{1}{(1+0.1t)^a}, 100(1+0.1t)^b)$, with $(a, b) = (0, 2)$, $(a, b) = (\frac{1}{2}, 1)$ and $(a, b) = (1, 0)$ (constant batch-sizes). Batch Sinkhorn algorithms uses $N = 100, 1000, 10000$. We train Sinkhorn on $t = 5000$ iterations, and train online Sinkhorn long enough to match the number of computations of the large Sinkhorn reference.

All OS convergence curves. To complete Fig. 1, Fig. 4 report the performance of online Sinkhorn for $\varepsilon \in \{10^{-4}, 10^{-3}, 10^{-2}, 10^{-1}\}$. The comparison of performance remains similar to the one produced in the main text.

Fully-corrective online Sinkhorn. Fig. 5 reports the performance of fully-corrected online Sinkhorn (FCOS). We observe that the fully-corrective scheme is less noisy than the non-corrected one. It is less efficient than OS on low-dimensional problems, but faster on the 10 dimensional problem. For GMM-10D, it outperforms the batch Sinkhorn algorithm with $N = 100, 1000$. Note

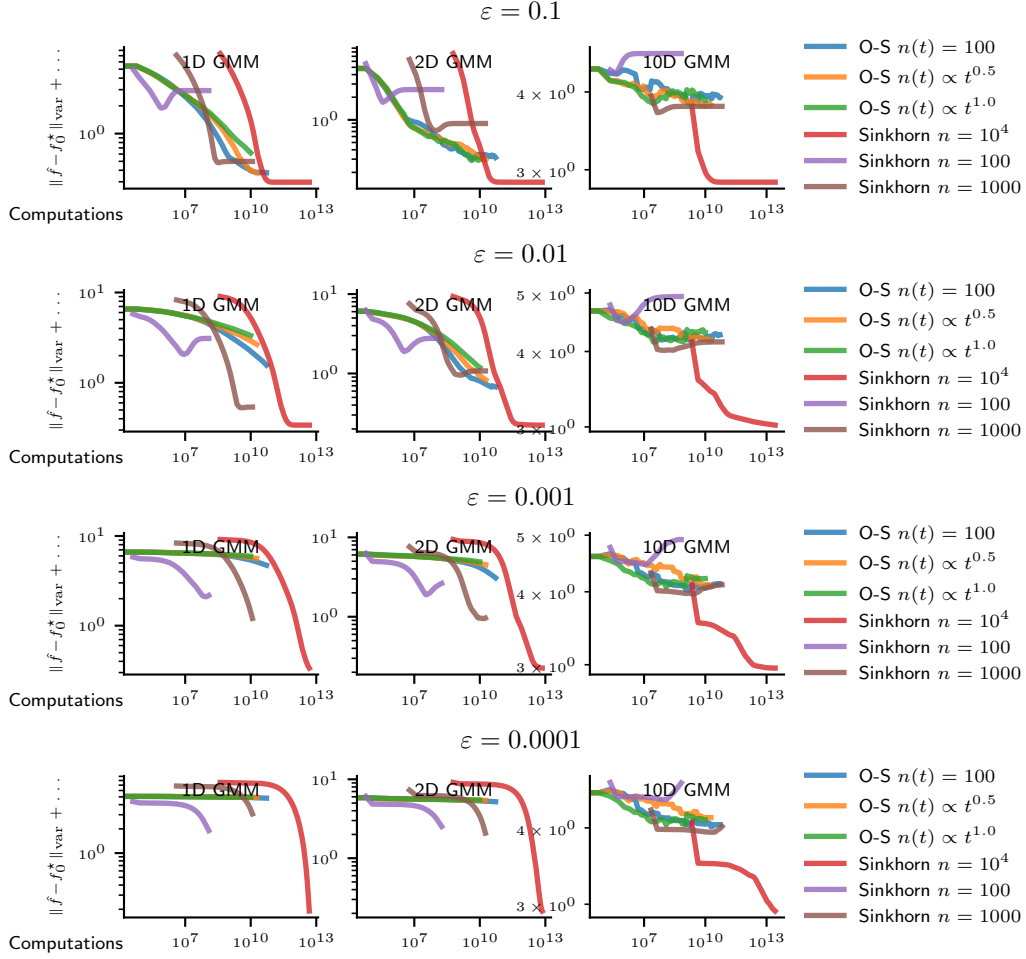


Figure 5: Performance of fully-corrective online Sinkhorn (O-S) for various ε .

that we interrupt FCOS for $n_t > 20,000$, as our implementation of the C -transform has a quadratic memory cost in n_t —this cost can be reduced to a linear cost with more careful implementation¹.

Randomized Sinkhorn. Fig. 6 reports the performance of randomized Sinkhorn. In low dimension, randomized Sinkhorn is a reasonable alternative to batch Sinkhorn, as it often outperforms it on average, for the same memory complexity (compare purple to orange curve for instance). In high dimension, batch Sinkhorn tend to perform slightly better.

C.2 OT between Gaussians

We measure the performance of online Sinkhorn to transport one Gaussian distribution α to another β . The potentials f^*, g^* are known exactly for this problem, which allows to have a strong golden standard. More precisely, adapting the formulae from [24], assuming $\alpha \sim \mathcal{N}(\mu, A)$ and $\beta \sim \mathcal{N}(\nu, B)$

¹Using e.g. <https://www.kernel-operations.io/keops/index.html>

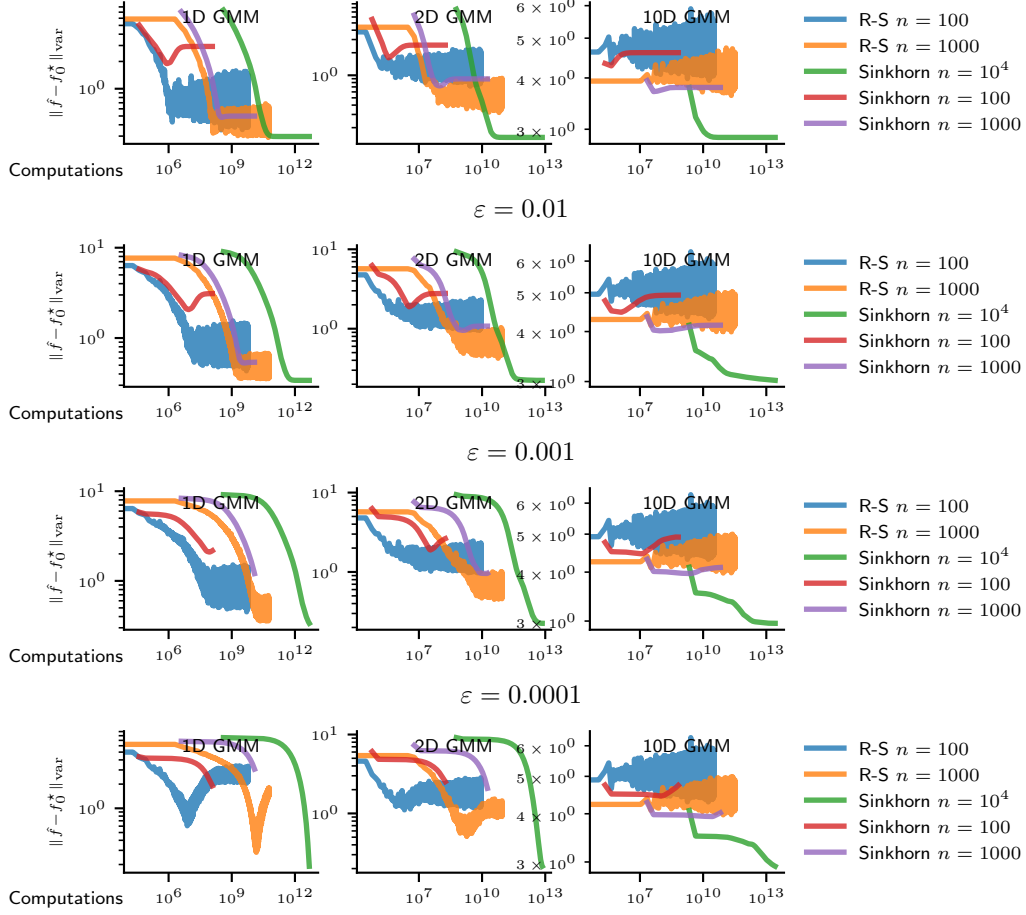


Figure 6: Performance of randomized Sinkhorn (R-S) for various ε .

and writing I the identity matrix in \mathbb{R}^d , we have

$$C \triangleq (AB + \frac{\varepsilon^2}{4}I)^{1/2}, \quad U \triangleq B(C + \frac{\varepsilon}{2}I)^{-1} - I, \quad V \triangleq A(C + \frac{\varepsilon}{2}I)^{-1} - I$$

$$f^* : x \rightarrow -\frac{1}{2}(x - \mu)^\top U(x - \mu) + x^\top(\mu - \nu)$$

$$g^* : y \rightarrow -\frac{1}{2}(y - \nu)^\top V(y - \nu) + y^\top(\nu - \mu)$$

We compare batch Sinkhorn ($N = 100, 1000, 10000$) to (non fully-corrected) online Sinkhorn, with $n(t) = B$, and $n(t) = B(1 + 0.1t)^{1/2}$, $B = 100$, and $\varepsilon \in \{10^{-4}, 10^{-3}, 10^{-2}, 10^{-1}\}$.

Results. As displayed in Fig. 7, online Sinkhorn outperforms batch Sinkhorn for all tested batch sizes and all ε . It is faster and does not converge towards biased potentials. This suggests that the performance of online Sinkhorn may be underestimated in the previous analyses due to poor potential reference.

C.3 Illustration of online Sinkhorn potentials on a 2D GMM

The estimate \hat{f}_t is useful to compute the gradient of the Sinkhorn distance $\mathcal{W}(\alpha, \beta)$ with respect to the distribution α . This is useful when α is a parametric distribution α_θ , as it allows to compute the gradient of the Sinkhorn distance with respect to θ using backpropagation. For simplicity, let us

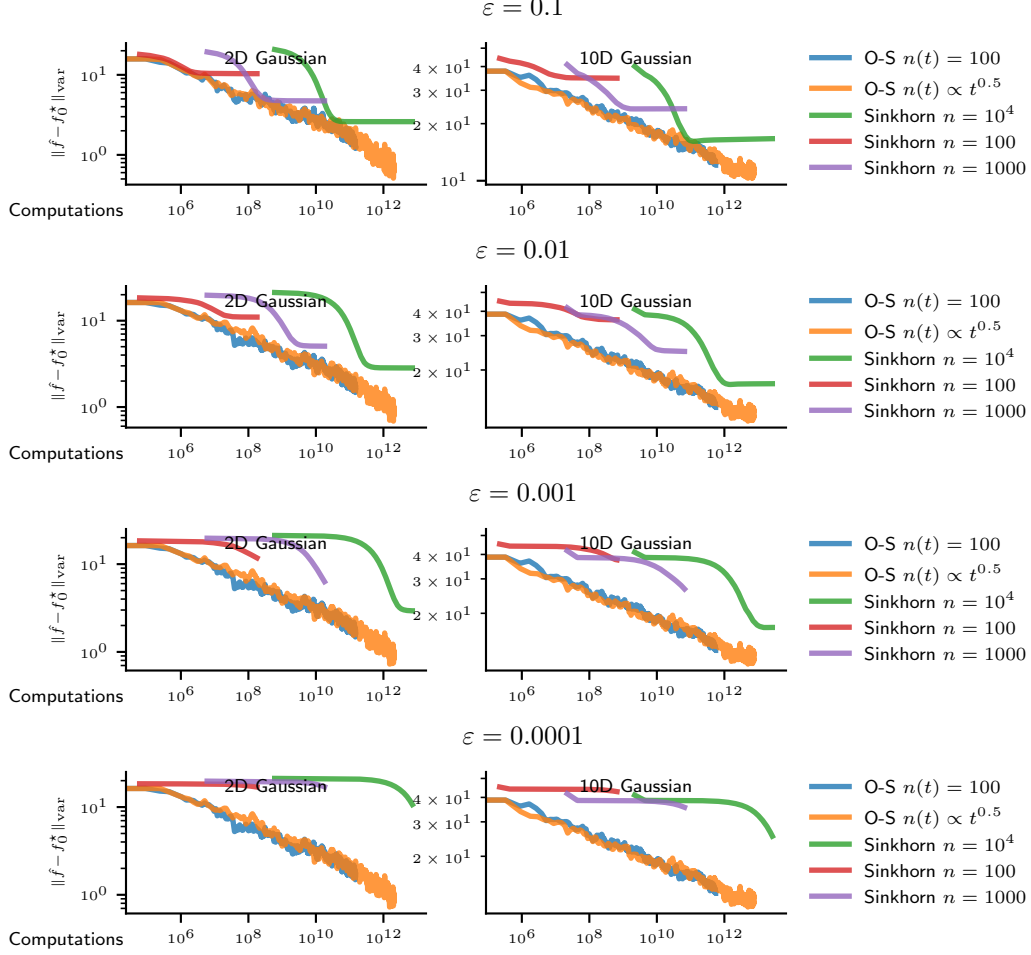


Figure 7: Performance of online-Sinkhorn to estimate OT between two Gaussians. Online Sinkhorn systematically outperforms batch Sinkhorn, but in term of speed and correction.

605 assume that $\alpha = \frac{1}{n} \sum_{i=1}^n \delta_{x_i}$. Then, for all $i \in [1, n]$,


$$\frac{\partial \mathcal{W}(\alpha, \beta)}{\partial x_i} = \nabla_x (x \rightarrow f^*(\alpha, \beta))(x_i),$$

606 so that $\nabla_x f^*(\alpha, \beta)$ provides a *displacement field* that can be descended to minimize $\alpha \rightarrow \mathcal{W}(\alpha, \beta)$.
 607 Such point of view can be extended to general distributions using the mean-field point of view, see
 608 e.g. [9, 36]. Estimating $\nabla_x f^*(\alpha, \beta)$ is therefore crucial to train e.g. generator networks. Both the
 609 online Sinkhorn and the batch Sinkhorn algorithm allow to estimate this vector field, through the
 610 plug-in estimator $x \rightarrow \nabla_x \hat{f}_t$, easily computed using the form (7) of \hat{f}_t .


611 **Experiment.** With 2D GMMs, we estimate a reference vector field ∇f_0^* using Sinkhorn on $N =$
 612 10,000 samples and qualitatively compare the estimations provided by online Sinkhorn and batch
 613 Sinkhorn ($N = 1,000$), for the same number of computations.

614 **Results.** We represent the estimations $\nabla_x \hat{f}_t$ in Fig. 8, for 10^8 computations. We compare them to a
 615 reference displacement field, estimated with 10^{10} computations. We observe that online Sinkhorn
 616 estimates a smoother displacement field than batch Sinkhorn for the same computational budget,
 617 that is closer to the reference displacement field. In particular, it is less noisy in low-mass areas.
 618 This suggest that online Sinkhorn would be a interesting replacement for batch Sinkhorn in training
 619 generative architectures (used by e.g. Genevay et al. [21]). α_θ is then defined as the push-forward of
 620 some simple measure with a neural network g_θ . We leave this direction for future work.


Online Sinkhorn $n(t) \propto 10t^{0.5}$, 10^8 computations



Sinkhorn $N=1000$, 10^8 computations



Sinkhorn $N=10000$, 10^{11} computations



Reg. OT displacement field
on empirical samples

Reg. OT displacement field
on a regular grid

Figure 8: Displacement field as defined by the potentials estimated by online-Sinkhorn and Sinkhorn on a 2D GMM. With the same computational budget, online Sinkhorn finds smoother displacement fields than Sinkhorn. Those are closer to the true reference displacement field (we use Sinkhorn on $N = 10000$ to estimate this reference). α and β log-likelihood level-lines are displayed in red and blue, while the arrows are proportional to $\nabla_x \hat{f}_t(x) d\alpha(x)$.

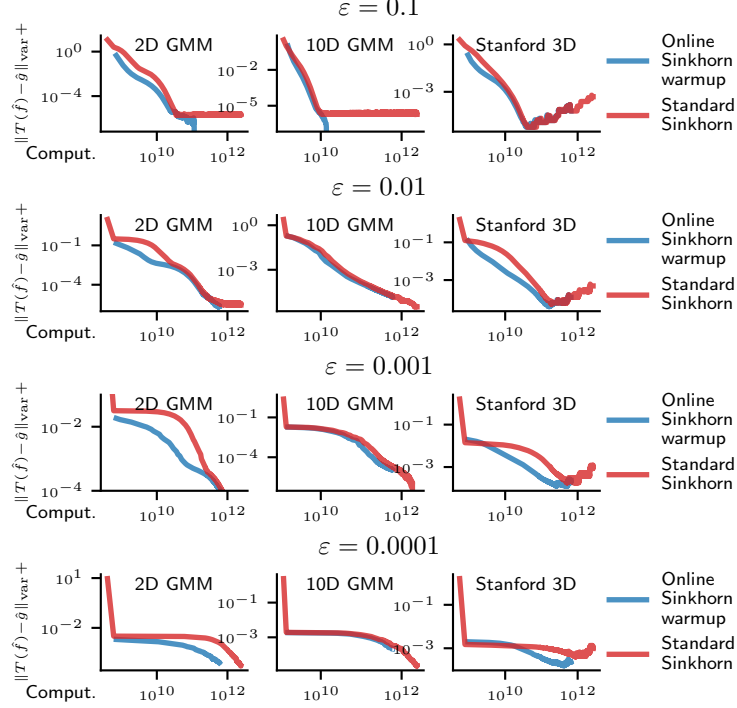


Figure 9: Performance of online-Sinkhorn as warmup for various ε .

621 C.4 Online Sinkhorn as a warmup process

622 **Grids and details for §5.2.** We set $(\eta_t, n(t)) = (\frac{1}{(1+0.1t)^a}, 100(1+0.1t)^b)$, with $(a, b) = (0, 2)$,
 623 $(a, b) = (\frac{1}{2}, 1)$ and $(a, b) = (1, 0)$ (constant batch-sizes). The batch Sinkhorn algorithm that is used for
 624 reference and after warmup uses $N = 10000$. In the reference algorithm, we precompute the distance
 625 matrix to save computation. In the warmup algorithm, this distance matrix is filled progressively and
 626 then kept in memory to perform C -transforms.

627 We evaluated OS and fully-corrective OS, and found that fully-corrective was less efficient (due to
 628 its higher cost in the early iterations). We evaluated sampling with and without replacement in the
 629 warmup phase, and found sampling without replacement to be more efficient.

630 **All warmup convergence curves.** To complete Fig. 3, we report convergence curves for different
 631 ε in Fig. 9. We find that speed-up increased with ε and both the 2D and 3D problems, but remains
 632 limited for the 10D problem.

D Stochastic mirror descent interpretation

The online Sinkhorn can be understood as a stochastic mirror descent algorithm for a non-convex problem. This equivalence is obtained by applying a change of variable in (1), defining

$$\mu \triangleq \alpha \exp(f) \quad \text{and} \quad \nu \triangleq \beta \exp(g). \quad (20)$$

The dual problem (2) rewrites as a minimisation problem over positive measures on \mathcal{X} and \mathcal{Y} :

$$-\min_{(\mu, \nu) \in \mathcal{M}^+(\mathcal{X})^2} \text{KL}(\alpha|\mu) + \text{KL}(\beta|\nu) + \langle \mu \otimes \nu, e^{-C} \rangle - 1, \quad (21)$$

where the function $\text{KL} : \mathcal{P}(\mathcal{X}) \times \mathcal{M}^+(\mathcal{X}) \triangleq \langle \alpha, \log \frac{d\alpha}{d\mu} \rangle$ is the Kullback-Leibler divergence between α and μ . This objective is block convex in μ, ν , but not jointly convex. As we now detail, this problem can be solved using a stochastic mirror descent [4], applied here over the Banach space of Radon measures on \mathcal{X} , equipped with the total variation norm.

Mirror maps and gradient. For this, we define the (convex) distance generating function $\mathcal{M}^+(\mathcal{X})^2 \rightarrow \mathbb{R}$:

$$\omega(\mu, \nu) \triangleq \text{KL}(\alpha|\mu) + \text{KL}(\beta|\nu).$$

The gradient of this function and of its Fenchel conjugate $\omega^* : \mathcal{C}(\mathcal{X})^2 \rightarrow \mathbb{R}$ yields two *mirror maps*. For all $(\mu, \nu) \in \mathcal{M}^+(\mathcal{X})^2$, $(\varrho, \varphi) \in \mathcal{C}(\mathcal{X})^2$, $\varrho < 0$, $\varphi < 0$,

$$\nabla \omega(\mu, \nu) = \left(-\frac{d\alpha}{d\mu}, -\frac{d\beta}{d\nu} \right) \quad \nabla \omega^*(\varrho, \varphi) = \left(-\frac{\alpha}{\varrho}, -\frac{\beta}{\varphi} \right).$$

The gradient $\nabla F(\mu, \nu)$ of the objective F appearing in (21) is a continuous function

$$\nabla_\mu F(\mu, \nu) = -\frac{1}{\frac{d\mu}{d\alpha}} + \int_{y \in \mathcal{X}} \frac{d\nu}{d\beta}(y) \exp(-C(\cdot, y)) d\beta(y)$$

and similarly for $\nabla_\nu F$.

Stochastic mirror descent. To define stochastic mirror descent iterations, we may replace integration over β by an integration over a sampled measure $\hat{\beta}$. This in turn defines an *unbiased gradient estimate* $\tilde{\nabla} F$ of ∇F , which has bounded second order moments. This absence of bias is crucial to prove convergence of SMD with high probability. Using the mirror maps and the stochastic estimation of the gradient, one has the following equivalence result, whose proofs stems from direct computations.

Proposition 7. *The stochastic mirror descent iterations*

$$(\mu_t, \nu_t) = \nabla \omega^* \left(\nabla \omega(\mu_t, \nu_t) - \eta_t \tilde{\nabla} F(\mu_t, \nu_t) \right)$$

are equal to the updates (6) under the change of variable (20).

Interpretation. It is important to realize that μ_t and ν_t do not need to be stored in memory. Instead, their associated potentials f_t and g_t are parametrized as (7). In particular, μ_t and ν_t remain absolutely continuous with respect to α and β respectively, so that the Kullback-Leibler divergence terms are always finite. Note that the mirror descent we consider operates in an infinite-dimensional space, as in Hsieh et al. [23].

Finally, we mention that when computing exact gradients (in the absence of noise) and when using constant step-size of $\eta_t = 1$, the algorithm matches exactly Sinkhorn iterations with simultaneous updates of the dual variables. This provides a novel interpretation on the Sinkhorn algorithm, that differs from the usual Bregman projection [5], and the related understanding of Sinkhorn as a constant step-size mirror descent on the primal objective [32] and on a semi-dual formulation [27].

Note that one can not directly apply the proofs of convergence of mirror descent to our problem, as the lack of convexity of problem (21) prevents their use.